

Evolutionary Dataset Optimisation: learning algorithm quality through evolution

Henry Wilde, Vincent Knight, Jonathan Gillard

Abstract

In this paper we propose a new method for learning how algorithms perform. Classically, algorithms are compared on a finite number of existing (or newly simulated) benchmark data sets based on some fixed metric. The algorithm(s) with the smallest value of this metric are chosen to be the ‘best performing’.

We offer a new approach to flip this paradigm. We instead aim to gain a richer picture of the performance of an algorithm by generating artificial data through genetic evolution, the purpose of which is to create populations of datasets for which a particular algorithm performs well. These data sets can be studied to learn as to what attributes lead to a particular progress of a given algorithm.

Following a detailed description of the algorithm as well as a brief description of an open source implementation, a number of numeric experiments are presented to show the performance of the method.

1 Introduction

This work presents a novel approach to learning the quality and performance of an algorithm through the use of evolution. When an algorithm is developed to solve a given problem, the designer is presented with questions about the performance of their proposed method, and its relative performance against existing methods. This is an inherently difficult task. However, under the current paradigm, the standard response to this situation is to use a known fixed set of datasets and a common metric amongst the proposed method and its competitors. The algorithm is then assessed based on this metric with minimal consideration for both the appropriateness or reliability of the datasets being used, and the robustness of the method in question.

This notion is not so easily observed when travelling in the opposite direction. Suppose that, instead, the benchmark was a dataset of particular interest and a preferable algorithm was to be determined for some task. There exist a number of methods employed across disciplines to complete this task that take into account the characteristics of the data and the context of the research problem. These methods include the use of diagnostic tests. For instance, in the case of clustering, if the data displayed an indeterminate number of non-convex blobs, then one could recommend that an appropriate clustering algorithm would be DBSCAN [5]. Otherwise, for scalability, k -means may be chosen.

The approach presented in this work aims to flip the paradigm described here by allowing the data itself to be unfixed. This fluidity in the data is achieved by generating data for which the algorithm performs well (or better than some other) through the use of an evolutionary algorithm. The purpose of doing so is not to simply create a bank of useful datasets but rather to allow for the subsequent studying of these datasets. In doing so, the attributes and characteristics which lead to the success (or failure) of the algorithm may be described, giving a broader understanding of the algorithm on the whole.

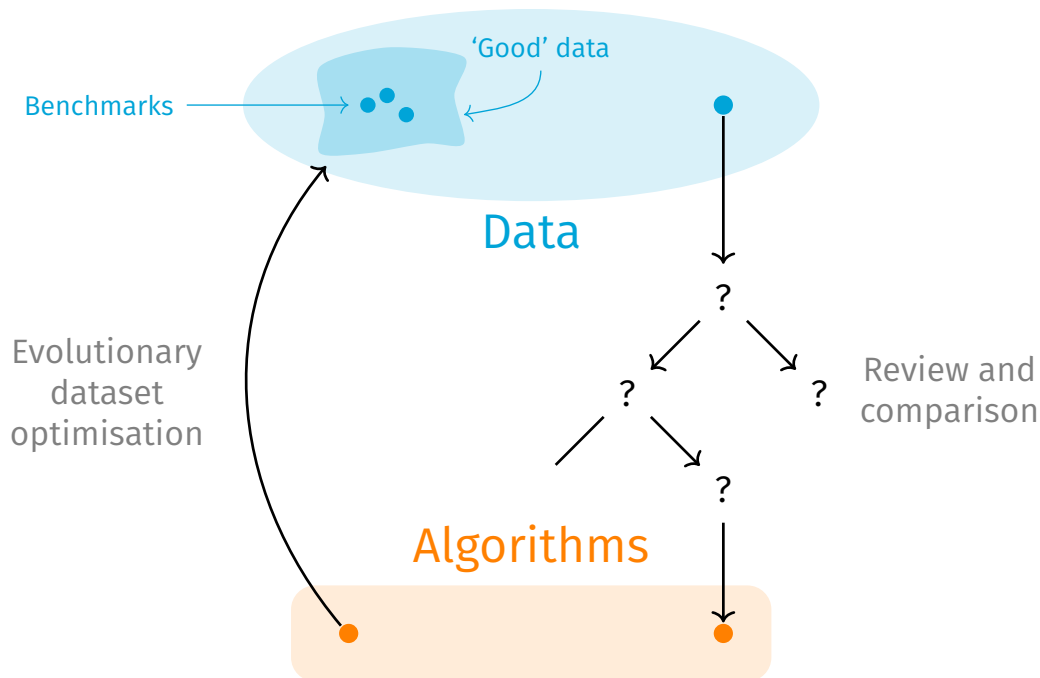


Figure 1: On the right: the current path for selecting some algorithm(s) based on their validity and performance for a given dataset. On the left: the proposed flip to better understand the space in which ‘good’ datasets exist for an algorithm.

This proposed flip has a number of motivations, and below is a non-exhaustive list of some of the problems that are presented by the established evaluation paradigm:

1. How are these benchmark examples selected? There is no true measure of their reliability other than their frequent use. In some domains and disciplines there are well-established benchmarks so those found through literature may well be reliable, but in others less so.
2. Sometimes, when there is a lack of benchmark examples, a ‘new’ dataset is simulated to assess the algorithm. This begs the question as to how and why that simulation is created. Not only this, but the origins of existing benchmarks is often a matter of convenience rather than their merit.
3. In disciplines where there are established benchmarks, there may still be underlying problems around the true performance of an algorithm:

- (i) As an example, work by Torralba and Efros [21] showed that image classifiers trained and evaluated on a particular dataset, or datasets, did not perform reliably when evaluated using other benchmark datasets that were determined to be similar. Thus leading to a model which lacks robustness.
- (ii) The amount of learning one can gain as to the characteristics of data which lead to good (or bad) performance of an algorithm is constrained to the finite set of attributes present in the benchmark data chosen in the first place.

Evolutionary algorithms (EAs) have been applied successfully to solve a wide array of problems - particularly where the complexity of the problem or its domain are significant. These methods are highly adaptive and their population-based construction (displayed in Figure 2) allows for the efficient solving of problems that are otherwise beyond the scope of traditional search and optimisation methods.

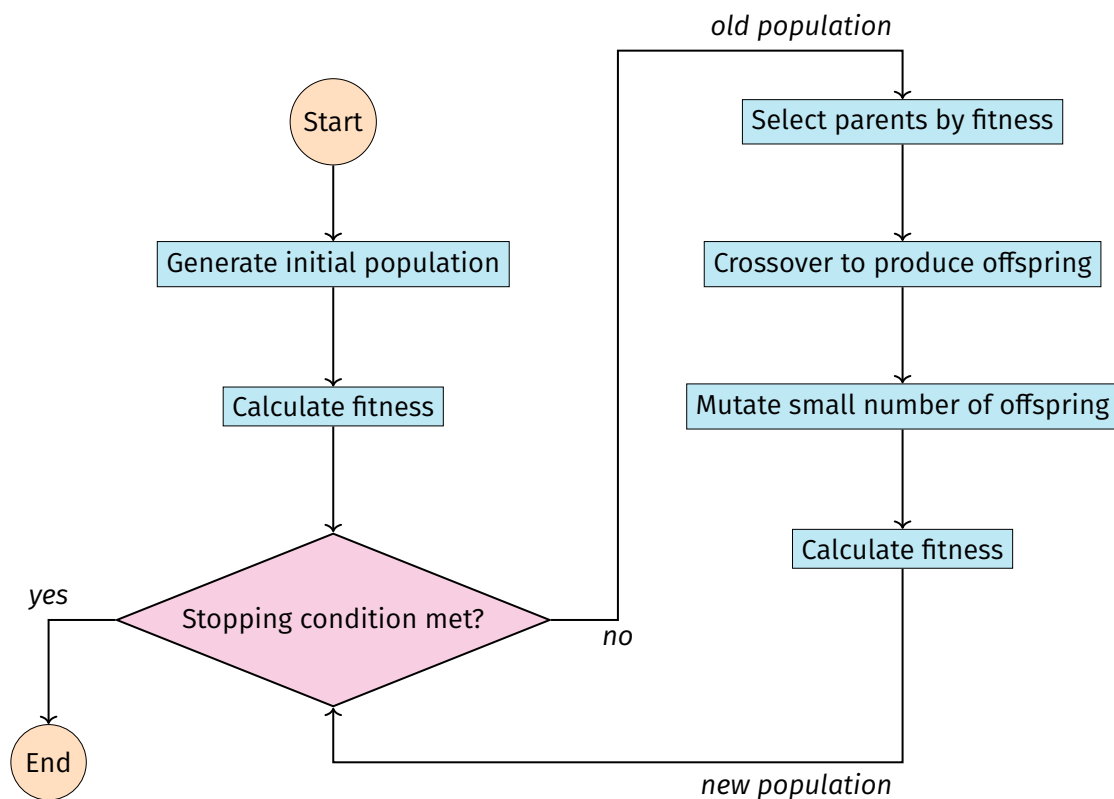


Figure 2: A general schematic for an evolutionary algorithm.

The use of EAs to generate artificial data is not a new concept. Its applications in data generation have included developing methods for the automated testing of software [8, 14, 18] and the synthesis of existing or confidential data [3]. Such methods also have a long history in the parameter optimisation of algorithms, and recently in the automated design of convolutional neural network (CNN) architecture [19, 20].

Other methods for the generation or synthesis of artificial data include simulated annealing [12] and generative adversarial networks (GANs) [6]. The unconstrained learning style of methods such as CNNs and GANs aligns with that proposed in this work. By allowing the EA to explore and learn about the search space in an organic way, less-prejudiced insight can be established that is not necessarily reliant on any particular framework or agenda.

Note that the proposed methodology is not simply to use an EA to optimise an algorithm over a search space with fixed dimension or datatype such as those set out in [3]. The size and sample space itself is considered as a property that can be traversed through the algorithm.

2 The evolutionary algorithm

2.1 Structure

In this section, the details of an algorithm that generates data for which a given function or, equivalently, an algorithm which is well suited, is described. This algorithm is to be referred to as “Evolutionary Dataset Optimisation” (EDO).

The EDO method is built as an evolutionary algorithm which follows a traditional (generic) schema with some additional features that keep the objective of artificial data generation in mind. With that, there are a number of parameters that are passed to EDO; the typical parameters of an evolutionary algorithm are a fitness function, f , which maps from an individual to a real number, as well as a population size, N , a maximum number of iterations, M , a selection parameter, b , and a mutation probability, p_m . In addition to these, EDO takes the following parameters:

- Limits on the number of rows an individual dataset can have:

$$R \in \{(r_{\min}, r_{\max}) \in \mathbb{N}^2 \mid r_{\min} \leq r_{\max}\}$$

- Limits on the number of columns a dataset can have:

$$C := (C_1, \dots, C_{|\mathcal{P}|}) \text{ where } C_j \in \{(c_{\min}, c_{\max}) \in (\mathbb{N} \cup \{\infty\})^2 \mid c_{\min} \leq c_{\max}\}$$

for each $j = 1, \dots, |\mathcal{P}|$. That is, C defines the minimum and maximum number of columns a dataset may have from each distribution in \mathcal{P} .

- A set of probability distribution families, \mathcal{P} . Each family in this set has some parameter limits which form a part of the overall search space. For instance, the family of normal distributions, denoted by $N(\mu, \sigma^2)$, would have limits on values for the mean, μ , and the standard deviation, σ .
- A probability vector to sample distributions from \mathcal{P} , $w = (w_1, \dots, w_{|\mathcal{P}|})$.
- A second selection parameter, $l \in [0, 1]$, to allow for a small proportion of “lucky” individuals to be carried forward.
- A shrink factor, $s \in [0, 1]$, defining the relative size of a component of the search space to be retained after adjustment.

The concepts discussed in this section form the mechanisms of the evolutionary dataset optimisation algorithm. To use the algorithm practically, these components have been implemented in Python as a library built on the scientific Python stack [13, 16]. The library is fully tested and documented (at <https://edo.readthedocs.io>) and is freely available online under the MIT license [4]. The EDO implementation was developed to be consistent with the best practices of open source software development.

The statement of the EDO algorithm is presented here to lay out its general structure from a high level perspective. Lower level discussion is provided below where additional algorithms

Algorithm 1: The evolutionary dataset optimisation algorithm

Input: $f, N, R, C, \mathcal{P}, w, M, b, l, p_m, s$

Output: A full history of the populations and their fitnesses.

begin

 create initial population of individuals

 find fitness of each individual

 record population and its fitness

while *current iteration less than the maximum* **and** *stopping condition not met* **do**

 select parents based on fitness and selection proportions

 use parents to create new population through crossover and mutation

 find fitness of each individual

 update population and fitness histories

if *adjusting the mutation probability* **then**

 update mutation probability

end

if *using a shrink factor* **then**

 shrink the mutation space based on parents

end

end

end

Algorithm 2: Creating a new population

Input: parents, $N, R, C, \mathcal{P}, w, p_m$

Output: A new population of size N

begin

 add parents to the new population

while *the size of the new population is less than N* **do**

 sample two parents at random

 create an offspring by crossing over the two parents

 mutate the offspring according to the mutation probability

 add the mutated offspring to the population

end

end

for the individual creation, evolutionary operator and shrinkage processes are given along with diagrams (where appropriate).

Note that there are no defined processes for how to stop the algorithm or adjust the mutation probability, p_m . This is down to their relevance to a particular use case. Some examples include:

- Stopping when no improvement in the best fitness is found within some K consecutive iterations [10].
- Utilising global behaviours in fitness to indicate a stopping point [11].
- Regular decreasing in mutation probability across the available attributes [9].

2.2 Individuals

Evolutionary algorithms operate in an iterative process on populations of individuals that each represent a solution to the problem in question. In a genetic algorithm, an individual is a solution encoded as a bit string of, typically, fixed length and treated as a chromosome-like object to be manipulated. In EDO, as the objective is to generate datasets and explore the space in which datasets exist, there is no encoding. As such the distinction is made that EDO is an evolutionary algorithm. Instead of this encoding, the produced datasets are manipulated directly so that the biological operators can be designed and be interpreted in a more meaningful way.

As is seen in Figure 3, an individual's creation is defined by the generation of its columns. A set of instructions on how to sample new values for that column are recorded in the form of a probability distribution. These distributions are sampled and created from the families passed in \mathcal{P} .




Figure 3: An example of how an individual is first created.

A caveat to this: one should not assume that the columns are a reliable representative of the distribution associated with them, or vice versa. This is particularly true of “shorter” datasets with a small number of rows, whereas confidence in the pair could be given more liberally for “longer” datasets with a larger number of rows. In any case, appropriate methods should be employed to understand the structure and characteristics of the data produced before formal conclusions are made.

2.3 Selection

The selection operator describes the process by which individuals are chosen from the current population to generate the next. Almost always, the likelihood of an individual being selected is determined by their fitness. This is because the purpose of selection is to preserve favourable qualities and encourage some homogeneity within future generations [2].

Algorithm 3: Creating an individual

Input: R, C, \mathcal{P}, w

Output: An individual defined by a dataset and some metadata

begin

 sample a number of rows and columns

 create an empty dataset

for each column in the dataset **do**

 sample a distribution from \mathcal{P}

 create an instance of the distribution

 fill in the column by sampling from this instance

 record the instance in the metadata

end

end

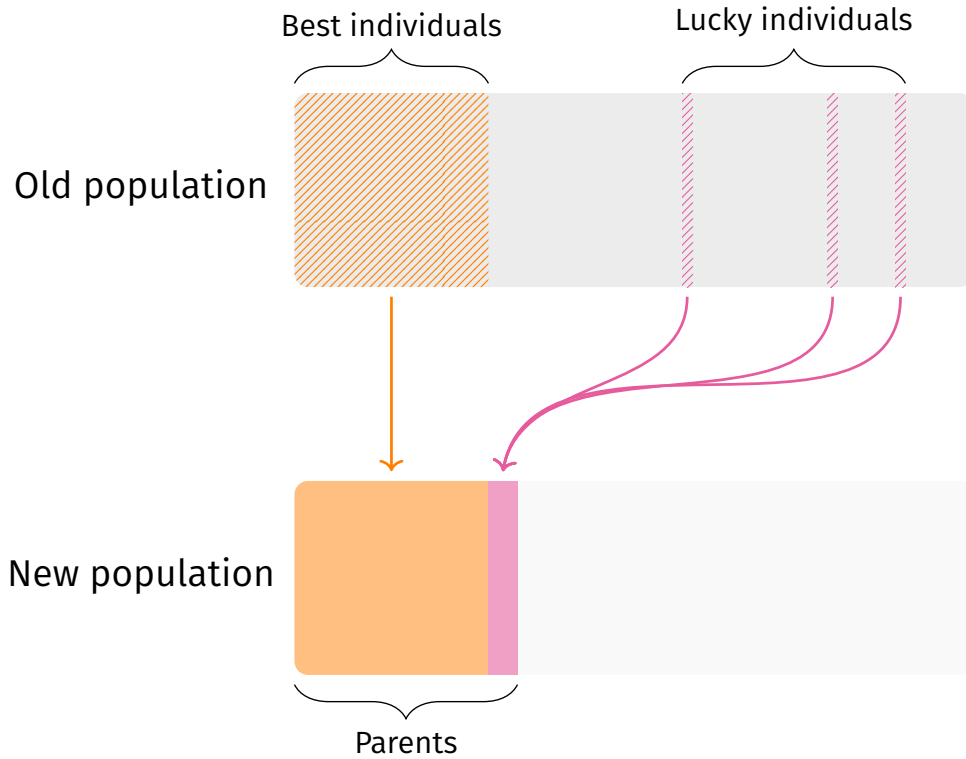


Figure 4: The selection process with the inclusion of some lucky individuals.

In EDO, a modified truncation selection method is used, as can be seen in Figure 4. Truncation selection takes a fixed number, $n_b = \lceil bN \rceil$, of the fittest individuals in a population and makes them the “parents” of the next. The modification is an optional stage after the best individuals

Algorithm 4: The selection process

Input: population, population fitness, b, l
Output: A set of parent individuals
begin
 calculate n_b and n_l
 sort the population by the fitness of its individuals
 take the first n_b individuals and make them parents
 if *there are any individuals left* **then**
 take the next n_l individuals and make them parents
 end
end

have been chosen: by passing some small l to the evolutionary algorithm, a number of the remaining individuals can be selected at random to be carried forward. This number is given by $n_l = \lceil lN \rceil$. It should be noted that even with this modification, no individual may be selected more than once in a single iteration but could potentially be present throughout the run of the algorithm.

It has been observed that, despite its efficiency as a selection operator, truncation selection can lead to premature convergence at local optima [7, 15]. Hence, allowing for the inclusion of a small number of randomly selected individuals may encourage diversity and exploration throughout the run of the algorithm.

A final component to this part of the evolutionary algorithm provides the ability to “shrink” the search space about the region observed in a given population. This method is based on a power law described in [1] that relies on a shrink factor, s . At each iteration, t , the parents are taken and every present distribution’s parameter limits, (l_t, u_t) , are centred about the mean parent value, μ :

$$l_{t+1} = \max \left\{ l_t, \mu - \frac{1}{2}(u_t - l_t)s^t \right\} \quad (1)$$

$$u_{t+1} = \max \left\{ u_t, \mu + \frac{1}{2}(u_t - l_t)s^t \right\} \quad (2)$$

This shrinking process is described in Algorithm 5 but note that the behaviour of this process

can give reductive results for some use cases. For instance, where there exist multiple optima for dataset characteristics.

Algorithm 5: Shrinking the mutation space

Input: parents, current iteration, \mathcal{P} , M , s

Output: A new mutation space focussed around the parents

begin

for *each distribution in \mathcal{P}* **do**

for *each parameter of the distribution* **do**

 get the current values for parameter over all parent columns

 find the mean of the current values

 find the new lower (1) and upper (2) bounds around the mean

 set the parameter limits

end

end

end

2.4 Crossover

Crossover is the operation of combining two individuals in order to create at least one offspring. In EDO, a method known as uniform crossover is used [17]. Under this method, a new individual is created by uniformly sampling each of its components from a set of two “parent” individuals. As can be seen in Figure 5, this method has been adapted for the dataset representation, i.e. two parent datasets have their dimensions and then columns sampled uniformly and without replacement to give a new individual.

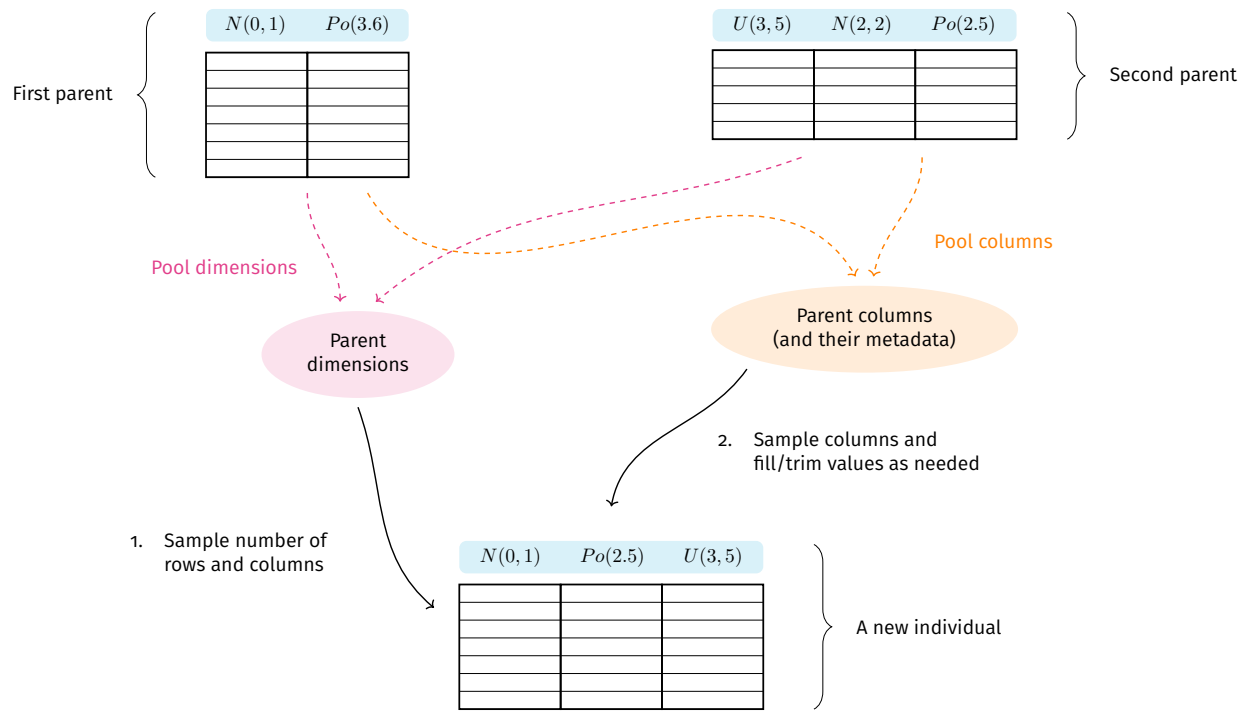


Figure 5: The crossover process.

Algorithm 6: The crossover process

Input: Two parents

Output: An offspring made from the parents ready for mutation

begin

 collate the columns and metadata from each parent in a pool

 sample each dimension from between the parents uniformly

 form an empty dataset with these dimensions

for *each column in the dataset* **do**

 sample a column (and its corresponding metadata) from the pool

if *this column is longer than required* **then**

 randomly select entries and delete them as needed

end

if *this column is shorter than required* **then**

 sample new values from the metadata and append them to the column as needed

end

 add this column to the dataset and record its metadata

end

end

2.5 Mutation

Mutation is used in evolutionary algorithms to encourage a broader exploration of the search space at each generation. Under this framework, the mutation process manipulates the phenotype of an individual where numerous things need to be modified including an individual's dimensions, column metadata and the entries themselves. This process is described in Figure 6.

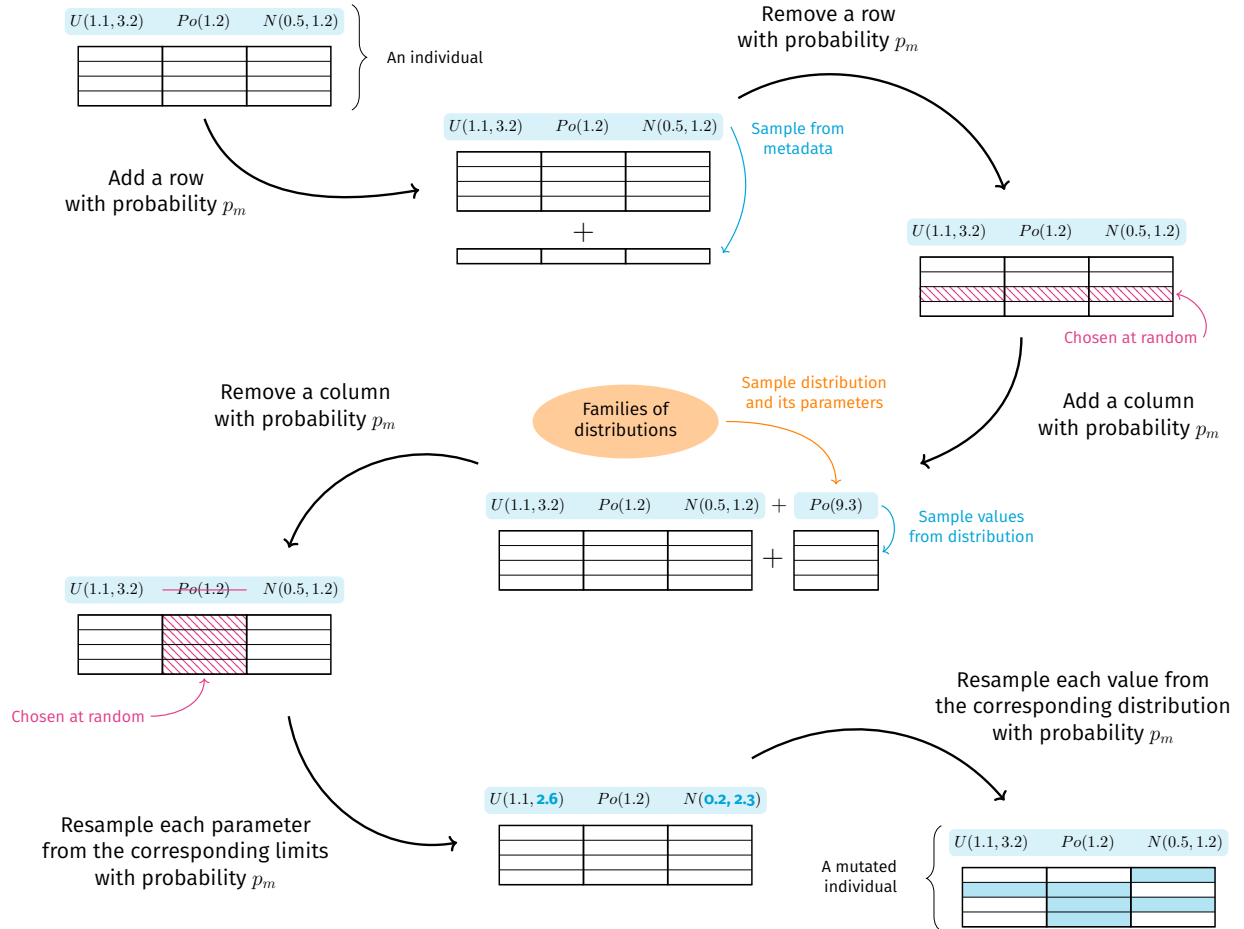


Figure 6: The mutation process.

As shown in Figure 6, each of the potential mutations occur with the same probability, p_m . However, the way in which columns are maintained assure that (assuming appropriate choices for f and \mathcal{P}) many mutations in the metadata and the dataset itself will only result in some incremental change in the individual's fitness relative to, say, a completely new individual.

Algorithm 7: The mutation process

Input: An individual, p_m , R , C , \mathcal{P} , w

Output: A mutated individual

begin

```
    sample a random number  $r \in [0, 1]$ 
    if  $r < p_m$  and adding a row would not violate  $R$  then
        sample a value from each distribution in the metadata
        append these values as a row to the end of the dataset
    end
    sample a new  $r \in [0, 1]$ 
    if  $r < p_m$  and removing a row would not violate  $R$  then
        remove a row at random from the dataset
    end
    sample a new  $r \in [0, 1]$ 
    if  $r < p_m$  and adding a new column would not violate  $C$  then
        create a new column using  $\mathcal{P}$  and  $w$ 
        append this column to the end of the dataset
    end
    sample a new  $r \in [0, 1]$ 
    if  $r < p_m$  and removing a column would not violate  $C$  then
        remove a column (and its associated metadata) at random from the dataset
    end
    for each distribution in the metadata do
        for each parameter of the distribution do
            sample a random number  $r \in [0, 1]$ 
            if  $r < p_m$  then
                sample a new value from within the distribution parameter limits
                update the parameter value with this new value
            end
        end
    end
    for each entry in the dataset do
        sample a random number  $r \in [0, 1]$ 
        if  $r < p_m$  then
            sample a new value from the associated column distribution
            update the entry with this new value
        end
    end
end
```

3 Numerical examples

- The x^2 example from the docs is a nice easy one to illustrate things
- Something stochastic
- Make use of the moving parts (linear focus of the mutation space, dwindling mutation probability, stopping conditions)
- If not the k -modes initialisation example, then another clustering one. Maybe k -means versus DBSCAN to show DBSCAN works better on non-convex clusters.

References

- [1] Adil Amirjanov. “Modeling the Dynamics of a Changing Range Genetic Algorithm”. In: *Procedia Computer Science* 102 (2016), pp. 570–577. DOI: <https://doi.org/10.1016/j.procs.2016.09.444>.
- [2] Thomas Bäck. “Selective pressure in evolutionary algorithms: a characterization of selection mechanisms”. In: *Proceedings of the First IEEE Conference on Evolutionary Computation. IEEE World Congress on Computational Intelligence*. 1994, pp. 57–62. DOI: 10.1109/ICEC.1994.350042.
- [3] Yingrui Chen, Mark Elliot, and Joseph Sakshaug. “A Genetic Algorithm Approach to Synthetic Data Production”. In: *PrAISe@ECAI*. 2016.
- [4] The EDO library developers. *EDO: v0.1*. 2018. DOI: 10.5281/zenodo.2557597. URL: <http://dx.doi.org/10.5281/zenodo.2557597>.
- [5] Martin Ester et al. “A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise”. In: *KDD*. 1996.
- [6] Ian Goodfellow et al. “Generative Adversarial Nets”. In: *Advances in Neural Information Processing Systems* 27. Curran Associates, Inc., 2014, pp. 2672–2680. URL: <http://papers.nips.cc/paper/5423-generative-adversarial-nets.pdf>.

- [7] Khalid Jebari. "Selection Methods for Genetic Algorithms". In: *International Journal of Emerging Sciences* 3 (Dec. 2013), pp. 333–344.
- [8] Chahine Koleejan, Bing Xue, and Mengjie Zhang. "Code coverage optimisation in genetic algorithms and particle swarm optimisation for automatic software test data generation". In: *2015 IEEE Congress on Evolutionary Computation (CEC)* (2015), pp. 1204–1211.
- [9] Matthias Kuehn, Thomas Severin, and Horst Salzwedel. "Variable Mutation Rate at Genetic Algorithms: Introduction of Chromosome Fitness in Connection with Multi-Chromosome Representation". In: *International Journal of Computer Applications* 72 (July 2013), pp. 31–38. DOI: 10.5120/12636-9343.
- [10] Yiu-Wing Leung and Yuping Wang. "An orthogonal genetic algorithm with quantization for global numerical optimization". In: *IEEE Transactions on Evolutionary Computation* 5.1 (2001), pp. 41–53. DOI: 10.1109/4235.910464.
- [11] Luis Mart et al. "A stopping criterion for multi-objective optimization evolutionary algorithms". In: *Information Sciences* 367–368 (2016), pp. 700 –718. DOI: <https://doi.org/10.1016/j.ins.2016.07.025>.
- [12] Justin Matejka and George Fitzmaurice. "Same Stats, Different Graphs: Generating Datasets with Varied Appearance and Identical Statistics Through Simulated Annealing". In: *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*. CHI '17. ACM, 2017, pp. 1290–1294. DOI: 10.1145/3025453.3025912.
- [13] Wes McKinney. *Data Structures for Statistical Computing in Python*. Proceedings of the 9th Python in Science Conference. [Online; accessed 2019-03-01]. 2010–. URL: <https://pandas.pydata.org/>.
- [14] Christoph C. Michael, Gary McGraw, and Michael Schatz. "Generating Software Test Data by Evolution". In: *IEEE Trans. Software Eng.* 27 (2001), pp. 1085–1110.
- [15] Tatsuya Motoki. "Calculating the Expected Loss of Diversity of Selection Schemes". In: *Evolutionary Computation* 10.4 (2002), pp. 397–422. DOI: 10.1162/106365602760972776.

- [16] Travis Oliphant. *NumPy: A guide to NumPy*. USA: Trelgol Publishing. [Online; accessed 2019-03-01]. 2006-. URL: <http://www.numpy.org/>.
- [17] Eugene Semenkin and Maria Semenkina. "Self-configuring Genetic Algorithm with Modified Uniform Crossover Operator". In: *Advances in Swarm Intelligence*. 2012, pp. 414–421. ISBN: 978-3-642-30976-2.
- [18] Hossein Sharifipour, Mojtaba Shakeri, and Hassan Haghighi. "Structural test data generation using a memetic ant colony optimization based on evolution strategies". In: *Swarm and Evolutionary Computation* 40 (2018), pp. 76 –91. DOI: <https://doi.org/10.1016/j.swevo.2017.12.009>.
- [19] Masanori Suganuma, Shinichi Shirakawa, and Tomoharu Nagao. "A genetic programming approach to designing convolutional neural network architectures". In: *GECCO*. 2017.
- [20] Yanan Sun et al. "Automatically Designing CNN Architectures Using Genetic Algorithm for Image Classification". In: *CoRR* abs/1808.03818 (2018).
- [21] A. Torralba and A. A. Efros. *Unbiased Look at Dataset Bias*. Proceedings of the 2011 IEEE Conference on Computer Vision and Pattern Recognition. 2011. DOI: 10.1109/CVPR.2011.5995347. URL: <http://dx.doi.org/10.1109/CVPR.2011.5995347>.