

Evolutionary Dataset Optimisation: learning algorithm quality through evolution

Henry Wilde · Vincent Knight · Jonathan Gillard

Received: date / Accepted: date

Abstract In this paper we propose a new method for learning how algorithms perform. Classically, algorithms are compared on a finite number of existing (or newly simulated) benchmark data sets based on some fixed metrics. The algorithm(s) with the smallest value of this metric are chosen to be the ‘best performing’.

We offer a new approach to flip this paradigm. We instead aim to gain a richer picture of the performance of an algorithm by generating artificial data through genetic evolution, the purpose of which is to create populations of datasets for which a particular algorithm performs well. These data sets can be studied to learn as to what attributes lead to a particular progress of a given algorithm.

Following a detailed description of the algorithm as well as a brief description of an open source implementation, a number of numeric experiments are presented to show the performance of the method which we call Evolutionary Dataset Optimisation.

Keywords Evolutionary algorithm · Optimisation · Algorithm design · Artificial data generation

1 Introduction

This work presents a novel approach to learning the quality and performance of an algorithm through the use of evolution. When an algorithm is developed to solve a given problem, the designer is presented with questions about the performance of their proposed method, and its relative performance against existing methods. This is an inherently difficult task. However, under the current paradigm, the standard response to this situation is to use a known fixed set of datasets - or simulate new data sets themselves - and a common metric amongst the proposed method and its competitors. The algorithm is then assessed based on this metric with often minimal

consideration for both the appropriateness or reliability of the datasets being used, and the robustness of the method in question.

This notion is not so easily observed when travelling in the opposite direction. Suppose that, instead, the benchmark was a dataset of particular interest and a preferable algorithm was to be determined for some task. There exist a number of methods employed across disciplines to complete this task that take into account the characteristics of the data and the context of the research problem. These methods include the use of diagnostic tests. For instance, in the case of clustering, if the data displayed an indeterminate number of non-convex blobs, then one could recommend that an appropriate clustering algorithm would be DBSCAN [6]. Otherwise, for scalability, k -means may be chosen [28].

The approach presented in this work aims to flip the paradigm described here by allowing the data itself to be unfixed. This fluidity in the data is achieved by generating data for which the algorithm performs well (or better than some other) through the use of an evolutionary algorithm. The purpose of doing so is not to simply create a bank of useful datasets but rather to allow for the subsequent studying of these datasets. In doing so, the attributes and characteristics which lead to the success (or failure) of the algorithm may be described, giving a broader understanding of the algorithm on the whole. Our framework is described in Figure 1.




Figure 1: On the right: the current path for selecting some algorithm(s) based on their validity and performance for a given dataset. On the left: the proposed flip to better understand the space in which ‘good’ datasets exist for an algorithm.

This proposed flip has a number of motivations, and below is a non-exhaustive list of some of the problems that are presented by the established evaluation paradigm:

1. How are these benchmark examples selected? There is no true measure of their reliability other than their frequent use. In some domains and disciplines there are well-established benchmarks so those found through literature may well be reliable, but in others less so.
2. Sometimes, when there is a lack of benchmark examples, a ‘new’ dataset is simulated to assess the algorithm. This begs the question as to how and why that simulation is created. Not only this, but the origins of existing benchmarks is often a matter of convenience rather than their merit.
3. In disciplines where there are established benchmarks, there may still be underlying problems around the true performance of an algorithm:
 - (i) As an example, work by Torralba and Efros [27] showed that image classifiers trained and evaluated on a particular dataset, or datasets, did not perform reliably when evaluated using other benchmark datasets that were determined to be similar. Thus leading to a model which lacks robustness.
 - (ii) The amount of learning one can gain as to the characteristics of data which lead to good (or bad) performance of an algorithm is constrained to the finite set of attributes present in the benchmark data chosen in the first place.

Evolutionary algorithms (EAs) have been applied successfully to solve a wide array of problems - particularly where the complexity of the problem or its domain are significant. These methods are highly adaptive and their population-based construction (displayed in Figure 2) allows for the efficient solving of problems that are otherwise beyond the scope of traditional search and optimisation methods.

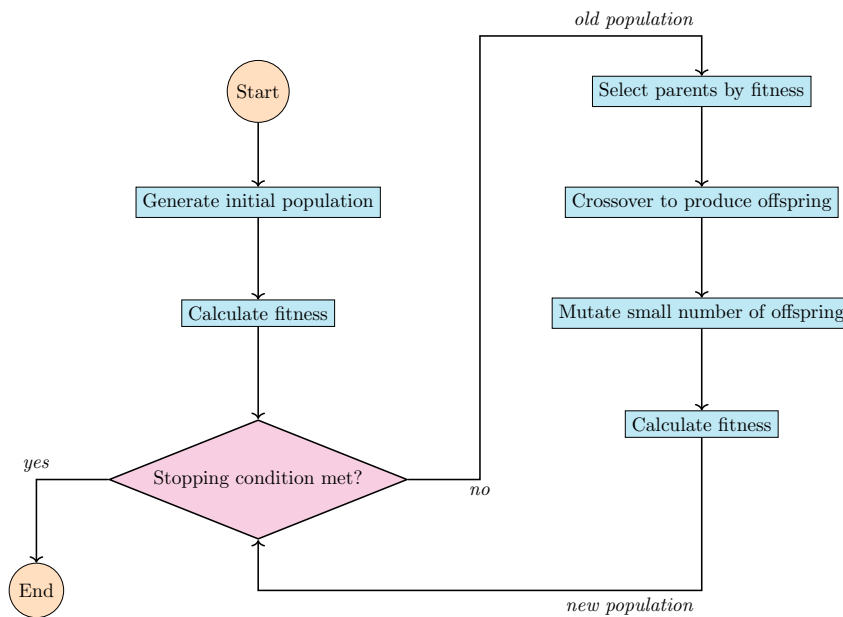


Figure 2: A general schematic for an evolutionary algorithm.

The use of EAs to generate artificial data is not a new concept. Its applications in data generation have included developing methods for the automated testing of software [11, 17, 22] and the synthesis of existing or confidential data [3]. Such methods also have a long history in the parameter optimisation of algorithms, and recently in the automated design of convolutional neural network (CNN) architecture [24, 25].

Other methods for the generation or synthesis of artificial data include simulated annealing [15] and generative adversarial networks (GANs) [7]. The unconstrained learning style of methods such as CNNs and GANs aligns with that proposed in this work. By allowing the EA to explore and learn about the search space in an organic way, less-prejudiced insight can be established that is not necessarily reliant on any particular framework or agenda.

Note that the proposed methodology is not simply to use an EA to optimise an algorithm over a search space with fixed dimension or datatype such as those set out in [3]. The size and sample space itself is considered as a property that can be traversed through the algorithm.

2 The evolutionary algorithm

In this section, the details of an algorithm that generates data for which a given function or, equivalently, an algorithm which is well suited, is described. This algorithm is to be referred to as “Evolutionary Dataset Optimisation” (EDO).

The EDO method is built as an evolutionary algorithm which follows a traditional (generic) schema with some additional features that keep the objective of artificial data generation in mind. With that, there are a number of parameters that are passed to EDO; the typical parameters of an evolutionary algorithm are a fitness function, f , which maps from an individual to a real number, as well as a population size, N , a maximum number of iterations, M , a selection parameter, b , and a mutation probability, p_m . In addition to these, EDO takes the following parameters:

- A set of probability distribution families, \mathcal{P} . Each family in this set has some parameter limits which form a part of the overall search space. For instance, the family of normal distributions, denoted by $N(\mu, \sigma^2)$, would have limits on values for the mean, μ , and the standard deviation, σ .
- A maximum number of “subtypes” for each family in \mathcal{P} . A subtype is an independent copy of the family that progresses separate from the others. These are the actual distribution objects which are traversed in the optimisation.
- A probability vector to sample distributions from \mathcal{P} , $w = (w_1, \dots, w_{|\mathcal{P}|})$.
- Limits on the number of rows an individual dataset can have,

$$R \in \{(r_{\min}, r_{\max}) \in \mathbb{N}^2 \mid r_{\min} \leq r_{\max}\}$$

- Limits on the number of columns a dataset can have,

$$C := (C_1, \dots, C_{|\mathcal{P}|}) \text{ where } C_j \in \{(c_{\min}, c_{\max}) \in (\mathbb{N} \cup \{\infty\})^2 \mid c_{\min} \leq c_{\max}\}$$

for each $j = 1, \dots, |\mathcal{P}|$. That is, C defines the minimum and maximum number of columns a dataset may have from each distribution in \mathcal{P} .

- A second selection parameter, $l \in [0, 1]$, to allow for a small proportion of ‘lucky’ individuals to be carried forward.
- A shrink factor, $s \in [0, 1]$, defining the relative size of a component of the search space to be retained after adjustment.

The concepts discussed in this section form the mechanisms of the evolutionary dataset optimisation algorithm. To use the algorithm practically, these components have been implemented in Python as a library built on the scientific Python stack [16, 19]. The library is fully tested and documented (at <https://edo.readthedocs.io>) and is freely available online under the MIT license [26]. The EDO implementation was developed to be consistent with the current best practices of open source software development [10].

Algorithm 1: The evolutionary dataset optimisation algorithm

Input: $f, N, R, C, \mathcal{P}, w, M, b, l, p_m, s$
Output: A full history of the populations and their fitnesses.
begin
 create initial population of individuals
 find fitness of each individual
 record population and its fitness
 while *current iteration less than the maximum and stopping condition not met* **do**
 select parents based on fitness and selection proportions
 use parents to create new population through crossover and mutation
 find fitness of each individual
 update population and fitness histories
 if *adjusting the mutation probability* **then**
 update mutation probability
 end
 if *using a shrink factor* **then**
 shrink the mutation space based on parents
 end
 end
end

Algorithm 2: Creating a new population

Input: parents, $N, R, C, \mathcal{P}, w, p_m$
Output: A new population of size N
begin
 add parents to the new population
 while *the size of the new population is less than N* **do**
 sample two parents at random
 create an offspring by crossing over the two parents
 mutate the offspring according to the mutation probability
 add the mutated offspring to the population
 end
end

The statement of the EDO algorithm is presented here to lay out its general structure from a high level perspective. Lower level discussion is provided below where additional algorithms for the individual creation, evolutionary operator and shrinkage processes are given along with diagrams (where appropriate). Note that there are no defined processes for how to stop the algorithm or adjust the mutation probability, p_m . This is down to their relevance to a particular use case. Some examples include:

- Regular decreasing in mutation probability across the available attributes [12].
- Stopping when no improvement in the best fitness is found within some K consecutive iterations [13].
- Utilising global behaviours in fitness to indicate a stopping point [14].

2.1 Individuals

Evolutionary algorithms operate in an iterative process on populations of individuals that each represent a solution to the problem in question. In a genetic algorithm, an individual is a solution encoded as a bit string of, typically, fixed length and treated as a chromosome-like object to be manipulated. In EDO, as the objective is to generate datasets and explore the space in which datasets exist, there is no encoding. As such the distinction is made that EDO is an evolutionary algorithm.

As is seen in Figure 3, an individual’s creation is defined by the generation of its columns. A set of instructions on how to sample new values (in mutation, for instance, Section 2.4) for that column are recorded in the form of a probability distribution. These distributions are sampled and created from the families passed in \mathcal{P} . In EDO, the produced datasets and their metadata are manipulated directly so that the biological operators can be designed and be interpreted in a more meaningful way as will be seen later in this section.

However, one should not assume that the columns are a reliable representative of the distribution associated with them, or vice versa. This is particularly true of ‘shorter’ datasets with a small number of rows, whereas confidence in the pair could be given more liberally for ‘longer’ datasets with a larger number of rows. In any case, appropriate methods for analysis should be employed before formal conclusions are made.

2.2 Selection

The selection operator describes the process by which individuals are chosen from the current population to generate the next. Almost always, the likelihood of an individual being selected is determined by their fitness. This is because the purpose of selection is to preserve favourable qualities and encourage some homogeneity within future generations [2].

In EDO, a modified truncation selection method is used [9], as can be seen in Figure 4. Truncation selection takes a fixed number, $n_b = \lceil bN \rceil$, of the fittest individuals in a population and makes them the ‘parents’ of the next. It has been observed that, despite its efficiency as a selection operator, truncation selection can lead to premature

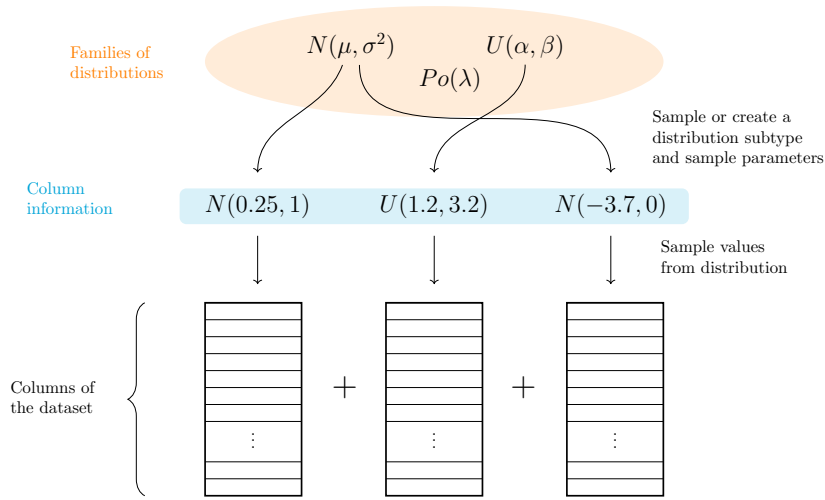


Figure 3: An example of how an individual is first created.

Algorithm 3: Creating an individual**Input:** R, C, \mathcal{P}, w **Output:** An individual defined by a dataset and some metadata**begin**

sample a number of rows and columns

create an empty dataset

for each column in the dataset do sample a distribution from \mathcal{P}

create an instance of the distribution

fill in the column by sampling from this instance

record the instance in the metadata

end**end****Algorithm 4:** The selection process**Input:** population, population fitness, b, l **Output:** A set of parent individuals**begin** calculate n_b and n_l

sort the population by the fitness of its individuals

 take the first n_b individuals and make them parents **if there are any individuals left then** take the next n_l individuals and make them parents **end****end**




Figure 4: The selection process with the inclusion of some lucky individuals.

convergence at local optima [9, 18]. The modification for EDO is an optional stage after the best individuals have been chosen: with some small l , a number, $n_l = \lceil lN \rceil$, of the remaining individuals can be selected at random to be carried forward. Hence, allowing for a small number of randomly selected individuals may encourage diversity and further exploration throughout the run of the algorithm. It should be noted that regardless of this step, an individual could potentially be present throughout the entirety of the algorithm.

After the parents have been selected, there are two adjustments made to the current search space. The first is that the subtypes for each family in \mathcal{P} are updated to only those present in the parents. The second adjustment is a process which acts on the distribution parameter limits for each subtype in \mathcal{P} . This adjustment gives the ability to ‘shrink’ the search space about the region observed in a given population. This method is based on a power law described in [1] that relies on a shrink factor, s . At each iteration, t , every distribution subtype which is present in the parents has its parameter’s limits, (l_t, u_t) , adjusted. This adjustment is such that the new limits, (l_{t+1}, u_{t+1}) are centred about the mean observed value, μ , for that parameter:

$$l_{t+1} = \max \left\{ l_t, \mu - \frac{1}{2}(u_t - l_t)s^t \right\} \quad (1)$$

$$u_{t+1} = \min \left\{ u_t, \mu + \frac{1}{2}(u_t - l_t)s^t \right\} \quad (2)$$

The shrinking process is given explicitly in Algorithm 5. Note that the behaviour of this process can produce reductive results for some use cases and is optional.

Algorithm 5: Shrinking the mutation space

Input: parents, current iteration, \mathcal{P}, M, s
Output: A new mutation space focussed around the parents

```

begin
  for each distribution subtype in  $\mathcal{P}$  do
    for each parameter of the distribution do
      get the current values for parameter over all parent columns
      find the mean of the current values
      find the new lower (1) and upper (2) bounds around the mean
      set the parameter limits
    end
  end
end

```

2.3 Crossover

Crossover is the operation of combining two individuals in order to create at least one offspring. In genetic algorithms, the term ‘crossover’ can be taken literally: two bit strings are crossed at a point to create two new bit strings. Another popular method is uniform crossover, which has been favoured for its efficiency and efficacy in combining individuals [21]. For EDO, this method is adapted to support dataset manipulation: a new individual is created by uniformly sampling each of its components (dimensions and then columns) from a set of two ‘parent’ individuals, as shown in Figure 5.

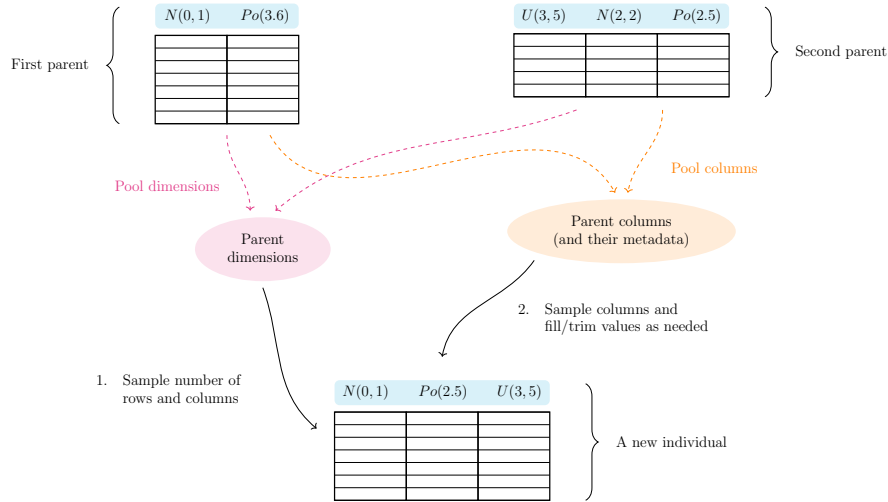


Figure 5: The crossover process between two individuals with different dimensions.

Observe that there is no requirement on the dimensions of the parents to be of similar or equal shapes. This is because the driving aim of the proposed method is to explore the space of all possible datasets. In the case where there is incongruence in the lengths of the two parents, missing values may appear in a shorter column that is sampled. To resolve this, values are sampled from the probability distribution associated with that column to fill in these gaps.

Algorithm 6: The crossover process

```

Input: Two parents
Output: An offspring made from the parents ready for mutation
begin
  collate the columns and metadata from each parent in a pool
  sample each dimension from between the parents uniformly
  form an empty dataset with these dimensions
  for each column in the dataset do
    sample a column (and its corresponding metadata) from the pool
    if this column is longer than required then
      | randomly select entries and delete them as needed
    end
    if this column is shorter than required then
      | sample new values from the metadata and append them to the column as needed
    end
    add this column to the dataset and record its metadata
  end
end

```

2.4 Mutation

Mutation is used in evolutionary algorithms to encourage a broader exploration of the search space at each generation. Under this framework, the mutation process manipulates the phenotype of an individual where numerous things need to be modified including an individual's dimensions, column metadata and the entries themselves. This process is described in Figure 6.

As shown in Figure 6, each of the potential mutations occur with the same probability, p_m . However, the way in which columns are maintained assure that (assuming appropriate choices for f and \mathcal{P}) many mutations in the metadata and the dataset itself will only result in some incremental change in the individual's fitness relative to, say, a completely new individual.

3 Examples

3.1 k -means clustering

The following examples act as a form of validation for EDO, and also highlight some of the nuances in its use. The objective of these examples is to use the proposed

Algorithm 7: The mutation process**Input:** An individual, $p_m, R, C, \mathcal{P}, w$ **Output:** A mutated individual**begin** sample a random number $r \in [0, 1]$ **if** $r < p_m$ *and adding a row would not violate R* **then**

sample a value from each distribution in the metadata

append these values as a row to the end of the dataset

end sample a new $r \in [0, 1]$ **if** $r < p_m$ *and removing a row would not violate R* **then**

remove a row at random from the dataset

end sample a new $r \in [0, 1]$ **if** $r < p_m$ *and adding a new column would not violate C* **then** create a new column using \mathcal{P} and w

append this column to the end of the dataset

end sample a new $r \in [0, 1]$ **if** $r < p_m$ *and removing a column would not violate C* **then**

remove a column (and its associated metadata) at random from the dataset

end **for each distribution in the metadata do** **for each parameter of the distribution do** sample a random number $r \in [0, 1]$ **if** $r < p_m$ **then**

sample a new value from within the distribution parameter limits

update the parameter value with this new value

end **end** **end** **for each entry in the dataset do** sample a random number $r \in [0, 1]$ **if** $r < p_m$ **then**

sample a new value from the associated column distribution

update the entry with this new value

end **end****end**

the only family in \mathcal{P} is:

$$\mathcal{U} := \{U(a, b) \mid a, b \in [0, 1]\} \quad (4)$$

The remaining parameters are as follows: $N = 100$, $R = (3, 100)$, $M = 1000$, $b = 0.2$, $l = 0$, $p_m = 0.01$, and shrinkage is excluded. Figure 7a shows an example of the fitness (above) and dimension (below) progression of the evolutionary algorithm under these conditions up until the 50th epoch.

There is a steep learning curve here; within the first 50 generations an individual is found with a fitness of roughly 10^{-10} which could not be improved on for a further 900 epochs. The same quick convergence is seen in the number of rows. This behaviour is quickly recognised as preferable and was dominant across all the trials conducted in this work. This preference for datasets with fewer rows is expected

given that I is the sum of the mean error from each cluster centre. With that, when k is fixed *a priori*, reducing the number of points in each cluster (i.e. the terms of the second summation) quickly reduces the mean error of that cluster and thus the value of I .




Figure 7a: Progressions for final inertia and dimension across the first 50 epochs with $R = (3, 100)$.

However, something that may be seen as unwanted is a compaction of the cluster centres. Referring to Figure 8a, the best and median individuals show two clusters that are essentially the same point whereas the worst is a random cloud across the whole of \mathcal{U} which was found in the initial population. The kind of behaviour exhibited by the best performing individuals here occurs in part because it is allowed. There are two immediate ways in which this is allowed: first, that a near-trivial case is included in R and, secondly, that the fitness function does nothing to penalise the proximity of the inter-cluster means, as well as aiming to reduce the intra-cluster means. This kind of unwanted behaviour highlights a subtlety in how EDO should be used; that experimentation and rigour are required to properly understand an algorithm's quality.

Hence, consider Figure 8b where the individuals have been generated with the same parameters as previously except with adjusted row limits, $R = (50, 100)$, so as to exclude this trivial case. In these trials, the results are equivalent: the worst performing individuals are without structure whilst the best-performing individuals display clusters that are dense about a single point despite the minimum number of rows being increased. Supposing this was not already a known result, we can




Figure 7b: Progressions for final inertia and dimension across the first 50 epochs with $R = (50, 100)$.

see mounting evidence in favour of this compaction being ‘optimal’ behaviour in a dataset for k -means clustering.

However, the fitness function may be addressed still, and more extensive studying may be done. Indeed, the final inertia could be considered a flawed or fragile fitness function if it is supposed to evaluate the efficacy of the k -means algorithm. Incorporating the inter-cluster spread to the fitness of an individual dataset would reduce this observed compaction. For instance, the silhouette coefficient is a metric used to evaluate the appropriateness of a clustering to a dataset and does precisely that. The silhouette coefficient of a clustering of a dataset is given by the mean of the silhouette value, $S(x)$, of each point $x \in Z_j$ in each cluster:

$$\begin{aligned}
 A(x) &:= \frac{1}{|Z_j| - 1} \sum_{y \in Z_j \setminus \{x\}} d(x, y), \\
 B(x) &:= \min_{k \neq j} \frac{1}{|Z_k|} \sum_{w \in Z_k} d(x, w), \\
 S(x) &:= \begin{cases} \frac{B(x) - A(x)}{\max\{A(x), B(x)\}} & \text{if } |Z_j| > 1 \\ 0 & \text{otherwise} \end{cases}
 \end{aligned} \tag{5}$$

The optimisation of the silhouette coefficient is analogous to finding a dataset which increases both the intra-cluster cohesion (the inverse of A) and inter-cluster

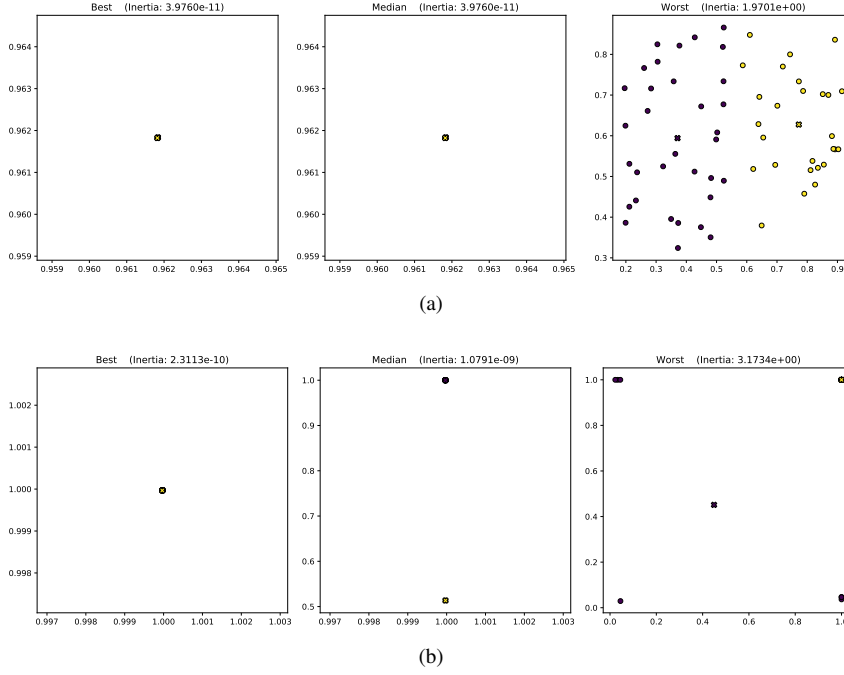


Figure 8: Representative individuals based on inertia with: (a) $R = (3, 100)$; (b) $R = (50, 100)$. Centroids displayed as crosses.

separation (B). Hence, the objective of minimising inertia is addressed by maximising cohesion. Meanwhile, the additional desire to spread out the clusters is considered by maximising separation.

Repeating the trials with the same parameters as with inertia, the silhouette fitness function yields the results summarised in Figures 9a and 9b. Irrespective of row limits, the datasets produced show increased separation from one another whilst maintaining low values in the final inertia of the clustering as shown in Figure 10. Again, the form of the individual clusters is much the same. The low values of inertia correspond to tight clusters, and the tightest clusters are those with a minimal number of points, i.e. a single point. As with the previous example, albeit at a much slower rate, the preferable individuals are those leading toward this case. That this gradual reduction in the dimension of the individuals occurs despite adjusting the fitness function and considering the space which excludes the trivial case bolsters the claim that the base case is also optimal.

Due to the nature of the implementation, any individual from any generation may be retrieved and studied should the final results be too concentrated on any given case. The summary provided here is one particular way of studying the body of datasets generated with this method and this transparency in the history and progression of the proposed method is something that sets it apart from other methods of the same ilk such as GANs which have a reputation of providing so-called ‘black box’ solutions.




Figure 9a: Progressions for silhouette and dimension across 1000 epochs at 100 epoch intervals with $R = (3, 100)$.




Figure 9b: Progressions for silhouette and dimension across 1000 epochs at 100 epoch intervals with $R = (50, 100)$.

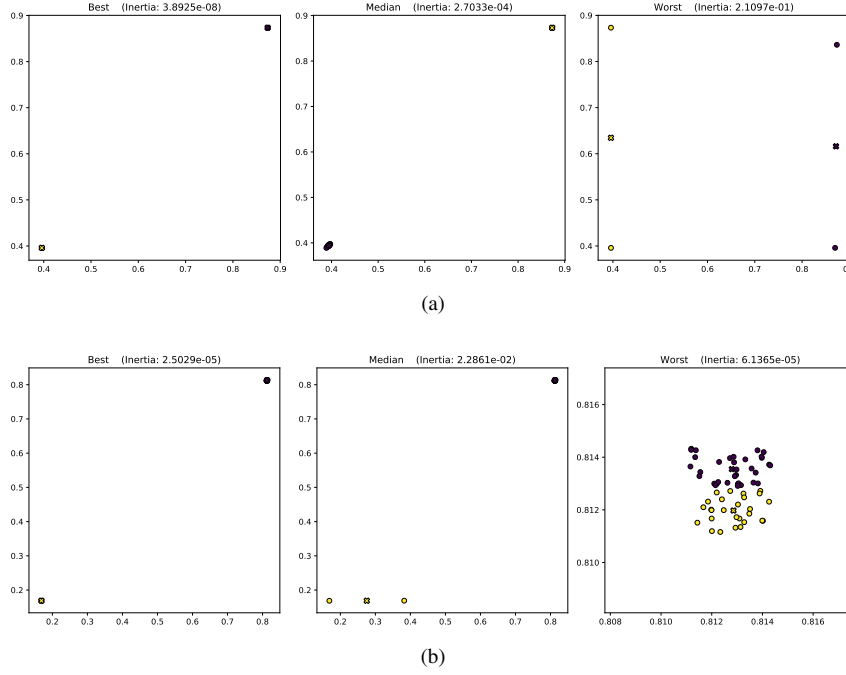


Figure 10: Representative individuals based on silhouette with: (a) $R = (3, 100)$; (b) $R = (50, 100)$. Centroids displayed as crosses.

3.2 Comparison with DBSCAN

The extent of the capabilities EDO holds as a tool to better understand an algorithm are especially apparent when comparing an algorithm against another (or set of others) simultaneously. This is done by utilising the freedom of choice in a fitness function for EDO. Consider two algorithms, A and B , and some common metric between them, g . Then their similarities and contrasts can be explored by considering the differences in this metric on the two algorithms. In terms of EDO, this means using $f = g_A - g_B$, $f = g_B - g_A$ or $f = |g_B - g_A|$ as the fitness function. By doing so, pitfalls, edge cases or fundamental conditions for the method may be highlighted. Overall, this process allows the researcher to more deeply learn about the method of interest beyond the traditional method of literature comparison on a particular example.

Consider the following use case with another clustering algorithm of a different form, Density Based Spatial Clustering of Applications with Noise (DBSCAN). In this particular case, the objective is to find datasets for which the method of interest, k -means, outperforms its alternative, DBSCAN. Here there is no concept of inertia as DBSCAN is density-based and is able to identify outliers [6]. As such, a valid metric must be chosen. One such metric is the silhouette score as defined in (5).

In this case, however, an adjustment to the fitness function must be made so as to accommodate for the condition of the silhouette coefficient that there must be more

than one cluster present. Let $S_k(X)$ and $S_D(X)$ denote the silhouette coefficients of the clustering found by k -means and DBSCAN respectively. Then the fitness function is defined to be:

$$f(X) = \begin{cases} S_D(X) - S_k(X), & \text{if DBSCAN identifies two or} \\ & \text{more clusters (inc. noise)} \\ \infty & \text{otherwise.} \end{cases} \quad (6)$$

There are several remarks to be made here. First, note the order of the subtraction here as EDO minimises fitness functions by default. Also, f takes values in the range $[-2, 2]$ where -2 is the best, i.e. $S_D(X) = -1$ and $S_k(X) = 1$. Likewise, 2 is the worst score. Finally, the silhouette coefficient requires at least two clusters to be present and so if DBSCAN identifies a single cluster then that individual will be penalised heavily under this fitness function when, in fact, that clustering may be of high quality. As such, this fitness function may require adjustment.

It must also be acknowledged that k -means and DBSCAN share no common parameters and so direct comparison is more difficult. For the purposes of this example, only one set of parameters is used but a thorough investigation should include a parameter sweep in similar, real-world use cases. The parameters being used are $k = 3$ for k -means, and $\epsilon = 0.1$, $MinPoints = 5$ for DBSCAN. This set was chosen following informal experimentation using the Python library Scikit-learn [20] to find comparable parameters in the given search space defined by the EDO parameters used previously with $R = (50, 100)$.

Figure 11 shows a summary of the progression of EDO for this use case. As with the previous examples where $R = (50, 100)$, the variation in the population fitness is unstable but there is a clear trend of improvement in the best individual over the course of the run. There is also a convergence seen in the number of rows a dataset has. The resting dimension varied across the trials conducted in this work but none exhibited a dramatic shift toward the lower limit of 50 rows as with previous examples. This is suggestive of a more competitive environment for individuals where slight changes to an individual can drastically alter their fitness.

The effect of such changes can be seen in Figure 12 where representative individuals are shown for this example. Here, the best performing individual, when clustered by k -means, shows three clear and nicely separated clusters. Note that they are not so tightly packed; again, this suggests that the route to an optimal individual is less clearly defined. In contrast, when the same dataset is clustered by DBSCAN a single cluster is found with a single noise point held within the convex hull of the cluster, i.e. there are overlapping clusters (since noise points form a single cluster). Hence, along with the fact that the larger cluster is widely spread, it follows that the clustering has a relatively small, negative silhouette coefficient.

Another point of interest here is the convexity of the clusters. A known condition for the success of k -means is that the presented clusters are of roughly equal size and are convex. This is due to the overall objective being to approximate the centroidal Voronoi tessellation [4]. Without this condition, up to the correct choice of k , the algorithm will fail to produce adequate results for either inertia or silhouette. DBSCAN, on the other hand, does not have this condition and is able to detect




Figure 11: Progressions for difference in silhouette (k -means-preferable) and dimension across 1000 epochs at 100 epoch intervals.

non-convex clusters so long as they are dense enough. Figure 12 shows the clustering found by each method and the respective convex and concave hulls of the clusters found. The ‘concave hull’ of a cluster is taken to be the α -shape of the cluster’s data points [5] where α is determined to be the smallest value such that all the points in the cluster are contained in a single polygon. The convexity of cluster Z_j , denoted \mathcal{C}_j , is then determined to be the ratio of the area of its concave hull, H_c , to the area of its convex hull, H_v [23]:

$$\mathcal{C}_j := \frac{\text{area}(H_c)}{\text{area}(H_v)} \quad (7)$$

With this definition, it should be clear that a perfectly convex cluster, such as a single point or line, would have $\mathcal{C}_j = 1$.

It can be seen that the convexity of the clustering found by k -means appears to be higher than that by DBSCAN. This was apparent across all trials conducted in this work and indicates that the condition for convex clusters is being sought out during the optimisation process. Meanwhile, however, it is not clear whether the performance of DBSCAN falls owing to its parameters or the method itself. This is a point where parameter sweeping would prove most useful so as to determine a crossing point for these two driving forces.

Now, to add to the discussion above, the inverse optimisation should be considered. That is, using the same parameters, the datasets for which DBSCAN outperforms k -means with respect to the silhouette coefficient are to be investigated. This is

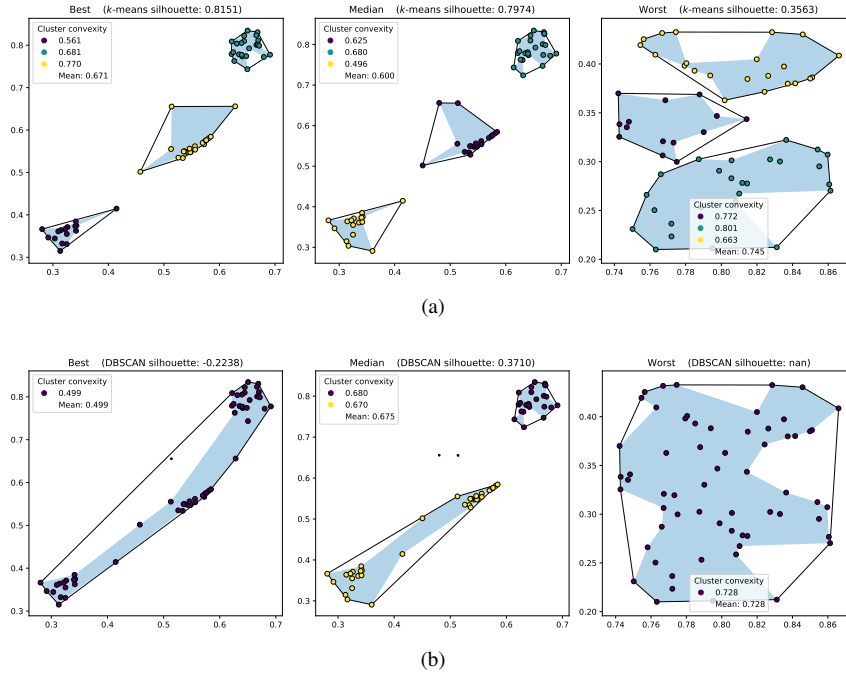


Figure 12: Representative individuals from a k -means-preferable run with clustering by: (a) k -means; (b) DBSCAN. Concave and convex hulls illustrated by shading and outline respectively.

equivalent to using $-f$ as the fitness function except with the same penalty of ∞ for the case set out in (6).

Figures 13 and 14 show the same summary as above with the revised fitness function. Inspecting the former, it is seen that the best fitness found is worse than with the previous example. This, in part, is due to the fact that k -means cannot find a clustering with negative values as no clusters may overlap. The method can, however, produce results with small silhouette scores where the clusters are tightly packed. Hence, the best fitness score is now -1 whereas the worst is 2, still.

Note in the first two frames of Figure 14a how k -means is forced to split what is evidently a single cluster in two whereas DBSCAN is able to identify the single cluster and the outlying noise (Figure 14b). The proximity of these clusters has then dragged the silhouette score down for k -means. Referring to Figure 14b, this kind of behaviour is certainly preferable for DBSCAN under these parameters: the beginning individuals are likely random clouds (as seen in the rightmost two frames of the figure) and the simplest step toward a fit dataset is one that maintains that vaguely dense body with minimal noise points far from it.




Figure 13: Progressions for difference in silhouette (DBSCAN-preferable) and dimension across 1000 epochs at 100 epoch intervals.

4 Conclusion

In this paper we have introduced a novel approach to understanding the quality of an algorithm by exploring the space in which their well-performing datasets exist. Following a detailed explanation of its internal mechanisms, a case study in k -means clustering was offered as validation for the proposed method. The method was able to reveal some known results without prior knowledge when investigating k -means in several scenarios, and again when comparing k -means and another leading clustering method, DBSCAN.

The method itself utilises biological operators to traverse the space of all possible datasets in an organic way with a minimal external framework attached. The generative nature of the proposed method also provides transparency and richness to the solution when compared to other contemporary techniques for artificial data generation as the entire history of individuals is preserved.

The evolutionary dataset optimisation method is dependent on a number of parameters set out in this paper and perhaps the most important of which is the choice of distribution families, \mathcal{P} ; these families set out the general statistical shape of the columns of the datasets that are produced and also control the present data types. The relationship between columns and their associated distribution is not causal and appropriate methods should be employed to understand the structure and characteristics of the data produced before formal conclusions are made as set out in the case study provided.

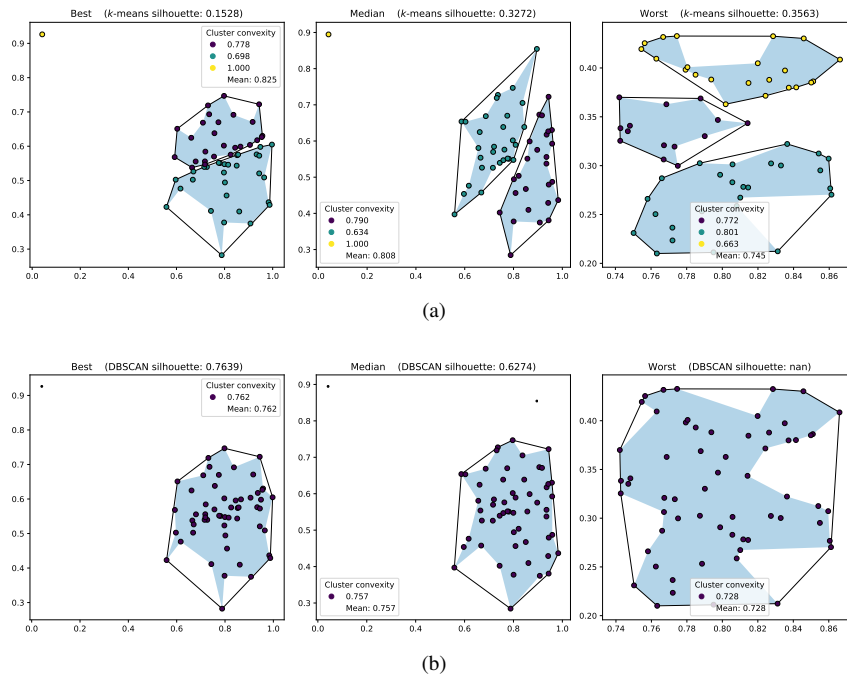


Figure 14: Representative individuals from a DBSCAN-preferable run with clustering by: (a) *k*-means; (b) DBSCAN. Concave and convex hulls illustrated by shading and outline respectively.

Acknowledgements The authors wish to thank the Cwm Taf Morgannwg University Health Board for their funding and support of the Ph.D. of which this work has formed a part.

Conflict of interest

The authors declare that they have no conflict of interest.

References

1. Amirjanov, A.: Modeling the dynamics of a changing range genetic algorithm. *Procedia Computer Science* **102**, 570 – 577 (2016). DOI <https://doi.org/10.1016/j.procs.2016.09.444>
2. Bäck, T.: Selective pressure in evolutionary algorithms: a characterization of selection mechanisms. In: *Proceedings of the First IEEE Conference on Evolutionary Computation. IEEE World Congress on Computational Intelligence*, pp. 57–62 (1994). DOI 10.1109/ICEC.1994.350042
3. Chen, Y., Elliot, M., Sakshaug, J.: A genetic algorithm approach to synthetic data production. In: *PrAISe@ECAI* (2016)
4. Du, Q., Emelianenko, M., Ju, L.: Convergence of the lloyd algorithm for computing centroidal voronoi tessellations. *SIAM Journal on Numerical Analysis* **44**(1), 102–119 (2006). DOI 10.1137/040617364
5. Edelsbrunner, H., Kirkpatrick, D., Seidel, R.: On the shape of a set of points in the plane. *IEEE Transactions on Information Theory* **29**(4), 551–559 (1983). DOI 10.1109/TIT.1983.1056714

6. Ester, M., Kriegl, H.P., Sander, J., Xu, X.: A density-based algorithm for discovering clusters in large spatial databases with noise. In: KDD (1996)
7. Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., Bengio, Y.: Generative adversarial nets. In: Advances in Neural Information Processing Systems 27, pp. 2672–2680. Curran Associates, Inc. (2014). URL <http://papers.nips.cc/paper/5423-generative-adversarial-nets.pdf>
8. Huang, Z.: Extensions to the k-means algorithm for clustering large data sets with categorical values. Data Mining and Knowledge Discovery **2**(3), 283–304 (1998). DOI 10.1023/A:1009769707641
9. Jebari, K.: Selection methods for genetic algorithms. International Journal of Emerging Sciences **3**, 333–344 (2013)
10. Jiménez, R.C., Kuzak, M., Alhamdoosh, M., Barker, M., Batut, B., Borg, M., Capella-Gutierrez, S., Chue Hong, N., Cook, M., Corpas, M., Flannery, M., Garcia, L., Gelpí, J.L., Gladman, S., Goble, C., González Ferreira, M., Gonzalez-Beltran, A., Griffin, P.C., Grüning, B., Hagberg, J., Holub, P., Hooft, R., Ison, J., Katz, D.S., Leskošek, B., López Gómez, F., Oliveira, L.J., Mellor, D., Mosbergen, R., Mulder, N., Perez-Riverol, Y., Pergl, R., Pichler, H., Pope, B., Sanz, F., Schneider, M.V., Stodden, V., Suchecki, R., SvobodováVařeková, R., Talvik, H.A., Todorov, I., Treloar, A., Tyagi, S., van Gompel, M., Vaughan, D., Via, A., Wang, X., Watson-Haigh, N.S., Crouch, S.: Four simple recommendations to encourage best practices in research software. F1000Research **6**, ELIXIR-876 (2017). DOI 10.12688/f1000research.11407.1
11. Koleejan, C., Xue, B., Zhang, M.: Code coverage optimisation in genetic algorithms and particle swarm optimisation for automatic software test data generation. 2015 IEEE Congress on Evolutionary Computation (CEC) pp. 1204–1211 (2015)
12. Kuehn, M., Severin, T., Salzwedel, H.: Variable mutation rate at genetic algorithms: Introduction of chromosome fitness in connection with multi-chromosome representation. International Journal of Computer Applications **72**, 31–38 (2013). DOI 10.5120/12636-9343
13. Leung, Y.W., Wang, Y.: An orthogonal genetic algorithm with quantization for global numerical optimization. IEEE Transactions on Evolutionary Computation **5**(1), 41–53 (2001). DOI 10.1109/4235.910464
14. Martí, L., García, J., Berlanga, A., Molina, J.M.: A stopping criterion for multi-objective optimization evolutionary algorithms. Information Sciences **367-368**, 700 – 718 (2016). DOI <https://doi.org/10.1016/j.ins.2016.07.025>
15. Matejka, J., Fitzmaurice, G.: Same stats, different graphs: Generating datasets with varied appearance and identical statistics through simulated annealing. In: Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems, CHI '17, pp. 1290–1294. ACM (2017). DOI 10.1145/3025453.3025912
16. McKinney, W.: Data structures for statistical computing in python. Proceedings of the 9th Python in Science Conference (2010–). URL <https://pandas.pydata.org/>. [Online; accessed 2019-03-01]
17. Michael, C.C., McGraw, G., Schatz, M.: Generating software test data by evolution. IEEE Trans. Software Eng. **27**, 1085–1110 (2001)
18. Motoki, T.: Calculating the expected loss of diversity of selection schemes. Evolutionary Computation **10**(4), 397–422 (2002). DOI 10.1162/106365602760972776
19. Oliphant, T.: NumPy: A guide to NumPy. USA: Trelgol Publishing (2006–). URL <http://www.numpy.org/>. [Online; accessed 2019-03-01]
20. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E.: Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research **12**, 2825–2830 (2011)
21. Semenkin, E., Semenkin, M.: Self-configuring genetic algorithm with modified uniform crossover operator. In: Advances in Swarm Intelligence, pp. 414–421 (2012)
22. Sharifipour, H., Shakeri, M., Haghighi, H.: Structural test data generation using a memetic ant colony optimization based on evolution strategies. Swarm and Evolutionary Computation **40**, 76 – 91 (2018). DOI <https://doi.org/10.1016/j.swevo.2017.12.009>
23. Sonka, M., Hlavac, V., Boyle, R.: Image Processing, Analysis and Machine Vision. Springer US (1993). DOI 10.1007/978-1-4899-3216-7
24. Suganuma, M., Shirakawa, S., Nagao, T.: A genetic programming approach to designing convolutional neural network architectures. In: GECCO (2017)
25. Sun, Y., Xue, B., Zhang, M., Yen, G.G.: Automatically designing CNN architectures using genetic algorithm for image classification. CoRR **abs/1808.03818** (2018)

26. The EDO library developers: EDO: v0.2.1 (2019). DOI 10.5281/zenodo.2651075. URL <http://dx.doi.org/10.5281/zenodo.2651075>
27. Torralba, A., Efros, A.A.: Unbiased look at dataset bias. Proceedings of the 2011 IEEE Conference on Computer Vision and Pattern Recognition (2011). DOI 10.1109/CVPR.2011.5995347. URL <http://dx.doi.org/10.1109/CVPR.2011.5995347>
28. Wu, X., Wu, X., Kumar, V.: The Top Ten Algorithms in Data Mining. CRC (2009)

A Appendix

A.1 Lloyd's algorithm

Algorithm 8: k -means (Lloyd's)

Input: a dataset X , a number of centroids k , a distance metric d
Output: a partition of X into k parts, Z

begin
 select k initial centroids, $z_1, \dots, z_k \in X$
 while *any point changes cluster or some stopping criterion is not met* **do**
 assign each point, $x \in X$, to cluster Z_{j^*} where:

$$j^* = \arg \min_{j=1, \dots, k} \left\{ d(x, z_j)^2 \right\}$$

 recalculate all centroids by taking the intra-cluster mean:

$$z_j = \frac{1}{|Z_j|} \sum_{x \in Z_j} x$$

end
end

A.2 Implementation example

Below is an example of how the Python implementation was used to complete the first example, including the definition of the fitness function.

```
import edo
from edo.pdfs import Uniform
from sklearn.cluster import KMeans

def fitness(dataframe, seed):
    """ Return the final inertia of 2-means on the 'dataframe'
    for the given 'seed'. """

    km = KMeans(n_clusters=2, random_state=seed).fit(dataframe)
    return km.inertia_

Uniform.param_limits["bounds"] = [0, 1]

pop_history, fit_history = edo.run_algorithm(
    fitness, size=100, row_limits=[3, 100], col_limits=[2, 2],
    families=[Uniform], max_iter=1000, best_prop=0.2,
    mutation_prob=0.01, seed=0, root="out",
    fitness_kwargs={"seed": 0},
)
```